

# Methods of Empirical Finance

Seminar (UE)

Christoph Huber  
University of Innsbruck

Master in Banking and Finance  
Winter term 2019/20 (this version: 2019-11-26)

# (Non)Parametric Statistical Tests

Methods of Empirical Finance

# (Non)Parametric Statistical Tests

## Parametric vs. Non-parametric tests

### Parametric:

- based on *parametric* families of *probability distributions*
- rely on assumptions that the data are drawn from a given prob. distribution (e.g. the underlying population is normally distributed)

### Non-parametric:

- *distribution-free* test statistics
- do *not* rely on assumptions about the data's prob. distribution

# (Non)Parametric Statistical Tests

## Parametric vs. Non-parametric tests

When should I use a parametric test?

When should I use a *non*-parametric test?

- **scale/level of measurement ?**  
non-metric (nominal, ordinal) vs. metric (interval, ratio) data
- **the distribution of the data ?**  
normally distributed vs. non-normally distributed

# (Non)Parametric Statistical Tests

## Parametric vs. Non-parametric tests

We mostly use ...

... **parametric** test statistics if ...

- the level of measurement is *metric* (interval or ratio),
- the population is *normally distributed*,
- the sample size is *large*

... **non-parametric** test statistics if ...

- the level of measurement is *non-metric* (nominal or ordinal),
- the population is *non-normally distributed*,
- the sample size is *small*

# (Non)Parametric Statistical Tests

As the appropriate test statistic depends on the data's level of measure, its (assumed) distribution, as well as its sample size, ...

... it makes sense to always get an overview about the data first,

... and — if you want to use parametric test statistics, in particular —, to test for normality of your data

# Testing for Normality


# Testing for Normality

## Some useful functions for distributions in R

For normal distributions:

```
dnorm(x=0, mean=0, sd=1)    # prob. density function (pdf)  
## [1] 0.3989423
```

```
x <- seq(-4, 4, length=100)  
y <- dnorm(x)  
plot(x, y)
```






# Testing for Normality

## Some useful functions for distributions in R

```
pnorm(74, mean=70, sd=2)    # cumulative density function
```

```
## [1] 0.9772499
```

```
y <- pnorm(x)  
plot(x, y)
```



# Testing for Normality

## Some useful functions for distributions in *R*

```
# e.g. find the Z-score of the 99th quantile of the standard normal distribution  
qnorm(.99, mean=0, sd=1)    # (inverse cumulative density function (cdf))
```

```
## [1] 2.326348
```

```
rnorm(5, mean = 10, sd = 2)    # generates normally distributed random variables
```

```
## [1] 12.441864 11.688925 7.233207 10.074942 11.343875
```

See [here](#), for example.


Similarly, these functions are available for most other popular probability distributions.

# Testing for Normality

## 'Descriptive': Histograms and normal density

```
hist(data$PCT_WOMEN_ON_BOARD, prob = TRUE, breaks = 20)  
curve(dnorm(x,  
          mean = mean(data$PCT_WOMEN_ON_BOARD, na.rm = TRUE),  
          sd = sd(data$PCT_WOMEN_ON_BOARD, na.rm = TRUE)),  
      col = "blue", lwd = 2,  
      add = TRUE, yaxt="n")
```


Histogram of data\$PCT\_WOMEN\_ON\_BOARD



# Testing for Normality

## 'Descriptive': QQ plots

```
qqnorm(data$PCT_WOMEN_ON_BOARD, frame = FALSE)  
qqline(data$PCT_WOMEN_ON_BOARD, col = ubluelight, lwd = 2)
```



# Testing for Normality

## Kolmogorov-Smirnov Goodness-of-Fit Test

tests whether the empirical distribution function of the sample and the cumulative distribution function of some reference distribution are equal in location and shape

- **nature of independent variable(s):** no independent variables, 1 population
- **scale of dependent variable:** ordinal or interval
- **hypothesis:**
  - $H_0: F(x) = F_0(x)$  for all values of  $x$
  - $H_1: F(x) \neq F_0(x)$  or  $F(x) > F_0(x)$  or  $F(x) < F_0(x)$

# Testing for Normality

## Kolmogorov-Smirnov Goodness-of-Fit Test — Example

- Test *whether the variable CHG\_PCT\_YTD is normally distributed*

```
ks.test(data$PCT_WOMEN_ON_BOARD,  
        "pnorm",  
        mean = mean(data$PCT_WOMEN_ON_BOARD, na.rm=TRUE),  
        sd = sd(data$PCT_WOMEN_ON_BOARD, na.rm=TRUE))
```

```
##  
##      One-sample Kolmogorov-Smirnov test  
##  
## data:  data$PCT_WOMEN_ON_BOARD  
## D = 0.075887, p-value = 0.006166  
## alternative hypothesis: two-sided
```

# Testing for Normality

## Shapiro-Wilk Test

tests whether a sample  $x_1, \dots, x_n$  is drawn from a normally distributed population

- **nature of independent variable(s):** no independent variables, 1 population
- **scale of dependent variable:** ordinal or interval
- **hypothesis:**
  - $H_0: F(x) = F_0(x)$  for all values of  $x$
  - $H_1: F(x) \neq F_0(x)$  or  $F(x) > F_0(x)$  or  $F(x) < F_0(x)$

# Testing for Normality

## Shapiro-Wilk Test — Example

- Test *whether the variable CHG\_PCT\_YTD is normally distributed*

```
shapiro.test(data$CHG_NET_YTD)
```

```
##  
##      Shapiro-Wilk normality test  
##  
## data:  data$CHG_NET_YTD  
## W = 0.37452, p-value < 2.2e-16
```



# One-Sample Test Statistics

# One-Sample Test Statistics

## $\chi^2$ Goodness of Fit Test

tests whether the observed proportions for a categorical variable of a single population differ significantly from hypothesized (expected) proportions

- **nature of independent variable(s):** no independent variables, 1 population
- **scale of dependent variable:** categorical/nominal
- **hypothesis:**
  - $H_0: p_i = p_{i,0} \forall i \in [1, \dots, k]$
  - $H_1: p_i \neq p_{i,0} \forall i \in [1, \dots, k]$  ('The observed proportions (  $p_{i,0}$  ) differ from hypothesized (expected) proportions (  $p_i$  ).')

# One-Sample Test Statistics

## $\chi^2$ Goodness of Fit Test — Example

- We use the data from `sp500_data.csv`.
- Test *whether the proportion of companies in the S&P 500 with a female CEO is different from 20%.*

```
females <- sum(data$FEMALE_CEO_OR_EQUIVALENT, na.rm=TRUE)
total <- sum(!is.na(data$FEMALE_CEO_OR_EQUIVALENT))
```

```
x <- c(total - females, females)
x
```

```
## [1] 479 23
```

```
chisq.test(x, p = c(0.8, 0.2))
```

```
##
##      Chi-squared test for given probabilities
##
## data:  x
## X-squared = 74.586, df = 1, p-value < 2.2e-16
```

# One-Sample Test Statistics

## One-Sample Wilcoxon Signed Rank Test

■ tests whether the population from which the data were sampled is symmetric or not about the default value

- **nature of independent variable(s):** no independent variables, 1 population
- **scale of dependent variable:** ordinal or interval

# One-Sample Test Statistics

## One-Sample Wilcoxon Signed Rank Test — Example

- Test *whether the median yearly net change in prices is different from 0*. I.e., whether the population from which the data were sampled is centered around 0.

```
wilcox.test(data$CHG_NET_YTD, mu = 0)
```

```
##  
##      Wilcoxon signed rank test with continuity correction  
##  
## data:  data$CHG_NET_YTD  
## V = 115647, p-value < 2.2e-16  
## alternative hypothesis: true location is not equal to 0
```

# One-Sample Test Statistics

## One-Sample $t$ -Test

evaluates whether the sample mean ( $\mu$ ) of a test variable significantly differs from a constant test value  $\mu_0$

- **nature of independent variable(s):** no independent variables, 1 population
- **scale of dependent variable:** interval or ratio
- **hypothesis:**
  - $H_0: \mu = \mu_0$
  - $H_1: \mu \neq \mu_0$  or  $\mu > \mu_0$  or  $\mu < \mu_0$  ('The sample mean  $\mu$  significantly differs from the constant test value  $\mu_0$ .')
- **assumptions:**
  - the observations in the sample are independent of each other
  - the observations follow a normal distribution with mean  $\mu$  and variance  $\sigma^2$

# One-Sample Test Statistics

## One-Sample $t$ -Test — Example

- Test *whether the mean yearly net change in prices is different from 0*.

```
summary(data$CHG_NET_YTD)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.     NA's  
## -146.950    5.395    15.170    28.076    33.490   1215.650         6
```

```
t.test(data$CHG_NET_YTD, mu = 0)
```

```
##  
##      One Sample t-test  
##  
## data:  data$CHG_NET_YTD  
## t = 9.1499, df = 498, p-value < 2.2e-16  
## alternative hypothesis: true mean is not equal to 0  
## 95 percent confidence interval:  
##  22.04771 34.10538  
## sample estimates:  
## mean of x  
##  28.07655
```

**But:** Do our assumptions hold? Is the data *normally distributed*?

# Unpaired Two-Sample Test Statistics



# Unpaired Two-Sample Test Statistics

## Pearson's $\chi^2$ Test for $r \times c$ tables

tests whether there is a relationship between two categorical variables (or tests whether or not the  $r$  samples are homogeneous regarding the proportion of observations in each of the  $c$  categories;  $\chi^2$  test for homogeneity)

- **nature of independent variable(s):** 1 independent variables with 2 levels (independent groups)
- **scale of dependent variable:** categorical/nominal
- **hypothesis:**
  - $H_0$ :  $O_{ij} = E_{ij}$  for all cells
  - $H_1$ :  $O_{ij} \neq E_{ij}$  for at least one cell

# Unpaired Two-Sample Test Statistics

## Pearson's $\chi^2$ Test for $r \times c$ tables — Example

- Test whether there is a relationship between the company's industry group and the percentage of women on its board.

```
chisq.test(data$INDUSTRY_GROUP, data$PCT_WOMEN_ON_BOARD)
```

```
##  
##      Pearson's Chi-squared test  
##  
## data:  data$INDUSTRY_GROUP and data$PCT_WOMEN_ON_BOARD  
## X-squared = 2532.7, df = 2419, p-value = 0.05262
```

# Unpaired Two-Sample Test Statistics

## Wilcoxon Rank Sum Test (Mann-Whitney $U$ -test)

■ tests whether two independent samples are drawn from the same population

- **nature of independent variable(s):** 1 independent variables with 2 levels (independent groups)
- **scale of dependent variable:** ordinal or interval

# Unpaired Two-Sample Test Statistics

## Wilcoxon Rank Sum Test — Example

- Test whether the percentage of women on the board of financial companies are drawn from the same distribution as the pct. of women on board of industrial companies.

```
wilcox.test(PCT_WOMEN_ON_BOARD ~ INDUSTRY_SECTOR,  
            data = data %>% filter(INDUSTRY_SECTOR == "Financial" | INDUSTRY_SECTOR == "Industrial"))
```

```
##  
##      Wilcoxon rank sum test with continuity correction  
##  
## data:  PCT_WOMEN_ON_BOARD by INDUSTRY_SECTOR  
## W = 3877, p-value = 0.04192  
## alternative hypothesis: true location shift is not equal to 0
```

# Unpaired Two-Sample Test Statistics

## Kolmogorov-Smirnov Test

tests whether two samples are drawn from the same distribution, or, more precisely, whether two independent empirical distributions are equal in location and shape

- **nature of independent variable(s):** 1 independent variables with 2 levels (independent groups)
- **scale of dependent variable:** ordinal or interval

# Unpaired Two-Sample Test Statistics

## Kolmogorov-Smirnov Test

- Test whether the percentage of women on the board of financial companies are drawn from the same distribution as the pct. of women on board of industrial companies or — more precisely — whether the the empirical distributions are equal in location and shape.

```
financial <- data %>% filter(INDUSTRY_SECTOR == "Financial") %>% .$PCT_WOMEN_ON_BOARD  
industrial <- data %>% filter(INDUSTRY_SECTOR == "Industrial") %>% .$PCT_WOMEN_ON_BOARD
```

```
ks.test(financial, industrial)
```

```
##  
##      Two-sample Kolmogorov-Smirnov test  
##  
## data:  financial and industrial  
## D = 0.23232, p-value = 0.02783  
## alternative hypothesis: two-sided
```

# Unpaired Two-Sample Test Statistics

## Independent Sample $t$ -Test

evaluates whether the sample means ( $\mu_1$  and  $\mu_2$ ) of two independent variables significantly differ

- **nature of independent variable(s):** 1 independent variables with 2 levels (independent groups)
- **scale of dependent variable:** interval or ratio
- **hypothesis:**
  - $H_0: \mu_1 = \mu_2$
  - $H_1: \mu_1 \neq \mu_2$  OR  $\mu_1 > \mu_2$  OR  $\mu_1 < \mu_2$
- **assumptions:**
  - all observations from both groups are independent of each other
  - the observations follow a normal distribution with mean  $\mu$  and variance  $\sigma^2$
  - the two independent samples have the same standard deviation, otherwise Welch's approximation should be applied

# Unpaired Two-Sample Test Statistics

## Independent Sample $t$ -Test — Example

- Test whether the mean percentage of woman is different in financial companies compared to industrial companies.

```
financial <- data %>% filter(INDUSTRY_SECTOR == "Financial") %>% .$PCT_WOMEN_ON_BOARD  
industrial <- data %>% filter(INDUSTRY_SECTOR == "Industrial") %>% .$PCT_WOMEN_ON_BOARD  
  
t.test(financial, industrial, var.equal=TRUE)
```

```
##  
##      Two Sample t-test  
##  
## data:  financial and industrial  
## t = 2.0421, df = 163, p-value = 0.04275  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##  0.07825703 4.65716398  
## sample estimates:  
## mean of x mean of y  
##  25.40932  23.04161
```



# Unpaired Two-Sample Test Statistics

## Independent Sample $t$ -Test — Example

- What if the two samples have different  $\sigma$ s?

```
t.test(financial, industrial, var.equal=FALSE)
```

```
##  
##      Welch Two Sample t-test  
##  
## data:  financial and industrial  
## t = 2.0669, df = 145.08, p-value = 0.04052  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##  0.1036396 4.6317814  
## sample estimates:  
## mean of x mean of y  
## 25.40932 23.04161
```

# Paired Two-Sample Test Statistics

# Paired Two-Sample Test Statistics

## Wilcoxon Signed Rank Test

tests whether two *dependent* samples represent two different populations or whether the two *dependent* samples are drawn from the same population

- **nature of independent variable(s):** 1 independent variables with 2 levels (*dependent* groups)
- **scale of dependent variable:** ordinal or interval
- **assumptions:**
  - the data are paired and come from the same population
  - each pair is chosen randomly and independently

# Paired Two-Sample Test Statistics

## Wilcoxon Signed Rank Test — Example

- Test whether the percentage of women on the board of financial companies in 2019 are drawn from the same distribution as the pct. of women on board of financial companies in 2014.

```
wilcox.test()
```

# Paired Two-Sample Test Statistics

## Dependent $t$ -Test for Paired Samples

evaluates whether the sample means ( $\mu_1$  and  $\mu_2$ ) of two *independent* samples differ significantly

- **nature of independent variable(s):** 1 independent variables with 2 levels (*dependent* groups)
- **scale of dependent variable:** interval or ratio
- **hypothesis:**
  - $H_0: \mu = \mu_0$
  - $H_1: \mu \neq \mu_0$  OR  $\mu > \mu_0$  OR  $\mu < \mu_0$  ('...')
- **assumptions:**
  - the observations follow a normal distribution with mean  $\mu$  and variance  $\sigma^2$
  - the two samples have the same standard deviation (homogeneity of variance)

# Paired Two-Sample Test Statistics

## Dependent $t$ -Test for Paired Samples

- Test whether the mean percentage of women on the board of financial companies in 2019 is different to the mean pct. of women on board of financial companies in 2014.

`t.test()`

# $n$ -Sample Generalizations: more than 2 groups

# $n$ -Sample Generalizations

## Kruskal-Wallis Test

tests whether two or more independent samples originate from the same distribution, i.e., extends the Mann-Whitney  $U$ -test when there are *more than two groups* (for two groups, both statistics are identical)

- **nature of independent variable(s):** 1 independent variables with two or more levels (independent groups)
- **scale of dependent variable:** ordinal or interval



# $n$ -Sample Generalizations

## Kruskal-Wallis Test — Example

- Test whether the percentage of women on the board of companies from different industry sectors are drawn from the same distribution.

```
kruskal.test(PCT_WOMEN_ON_BOARD ~ INDUSTRY_SECTOR,  
             data = data)
```

```
##  
##      Kruskal-Wallis rank sum test  
##  
## data:  PCT_WOMEN_ON_BOARD by INDUSTRY_SECTOR  
## Kruskal-Wallis chi-squared = 9.3157, df = 8, p-value = 0.3164
```

# $n$ -Sample Generalizations

## Jonkheere-Terpstra Test (Jonkheere's trend test)

tests whether three or more independent samples originate from the same distribution, (just as the Kruskal-Wallis test) but with an *ordered* alternative hypothesis

- **nature of independent variable(s):** 1 independent variables with three or more levels (independent groups)
- **scale of dependent variable:** ordinal or interval
- **hypothesis:**
  - $H_0: \theta_1 = \theta_2 = \dots = \theta_k$
  - $H_1: \theta_1 \leq \theta_2 \leq \dots \leq \theta_k$
- **assumptions and limitations:**
  - all observations from all  $k$  groups are independent of each other
  - hypothesized ordering of groups' rank-sums required

# $n$ -Sample Generalizations

## Jonkheere-Terpstra Test — Example

- Test whether the percentage of women on the board of companies from financial companies is larger than the percentage in industrial companies, which is then larger than the respective percentage in basic materials companies.

```
library(clinfun)
groups <- data %>% filter(INDUSTRY_SECTOR == "Financial" |
                        INDUSTRY_SECTOR == "Industrial" |
                        INDUSTRY_SECTOR == "Basic Materials")
groups$INDUSTRY_SECTOR <- factor(groups$INDUSTRY_SECTOR, levels = c("Financial", "Industrial", "Basic Materials"), ordered = TRUE)
jonckheere.test(groups$PCT_WOMEN_ON_BOARD,
                groups$INDUSTRY_SECTOR)
```

```
##
##      Jonckheere-Terpstra test
##
## data:
## JT = 4237, p-value = 0.1625
## alternative hypothesis: two.sided
```

# Permutation Tests

# Permutation Tests

Consider the following hypothetical problem:

- Ten subjects have been randomly assigned to one of two treatment conditions (A or B) and an outcome variable (score) has been recorded.

Treatment A	Treatment B
40	57
57	64
45	55
55	62
58	65

- First, let's use a parametric approach to compare Treatment A and Treatment B. → *t*-test:
  - assume normal distributions, independent groups, and equal variances

```
a <- c(40, 57, 45, 55, 58)
b <- c(57, 64, 55, 62, 65)
t.test(a, b, var.equal = TRUE)
```

```
##
##      Two Sample t-test
##
## data:  a and b
## t = -2.345, df = 8, p-value = 0.04705
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -19.0405455  -0.1594545
## sample estimates:
## mean of x mean of y
##      51.0      60.6
```

# Permutation Tests

A permutation test takes a different approach: If the two treatments are truly equivalent, the label (Treatment A or Treatment B) assigned to an observed score is arbitrary.

1. Calculate the observed t-statistic, as in the parametric approach; call this  $t_0$ .
2. Place all 10 scores in a single group.
3. Randomly assign five scores to Treatment A and five scores to Treatment B.
4. Calculate and record the new observed t-statistic.
5. Repeat steps 3–4 for every possible way of assigning five scores to Treatment A and five scores to Treatment B. There are 252 such possible arrangements.
6. Arrange the 252 t-statistics in ascending order. This is the empirical distribution, based on (or conditioned on) the sample data.
7. If  $t_0$  falls outside the middle 95 percent of the empirical distribution, reject the null hypothesis that the population means for the two treatment groups are equal at the 0.05 level of significance.

# Permutation Tests

- In both the permutation and the parametric approach, the exact same  $t$ -statistic is calculated.
- In the permutation approach, however, the statistic is compared not to a theoretical distribution (e.g. an assumed normal distribution) but to an *empirical distribution* created from *permutations of the observed data*.
- If this empirical distribution contains *all possible permutations* of the data, the permutation test is called an 'exact' test.
- If there are too many possible permutations to compute the test statistic in appropriate time, we can use (Monte Carlo) *simulations*.

# Permutation Tests

In *R*, we can use the `coin` package to run permutation tests as alternatives for a number of tests.

```
library(coin)
```

```
## Loading required package: survival
```

E.g.

```
df <- data.frame(value = c(a, b), treatment = c(rep("A", 5), rep("B", 5)))  
oneway_test(value ~ treatment, data = df)
```

```
##  
##      Asymptotic Two-Sample Fisher-Pitman Permutation Test  
##  
## data:  value by treatment (A, B)  
## Z = -1.9147, p-value = 0.05553  
## alternative hypothesis: true mu is not equal to 0
```

See [here](#), for example.



# (Non)parametric Correlation Measures

- Parametric: Pearson's  $\rho$
- Non-parametric:
  - Spearman rank correlation  $\rho$ : i.e., the Pearson correlation coefficient between the ranked variables
  - Kendall's  $\tau$ : similar to Spearman's  $\rho$ ; measures *ordinal association* between two measured quantities

The argumentations for using non-parametric *or* parametric correlation measures are very similar to the ones for using non\_parametric or parametric test statistics.

Things to consider are: level of measurement, assumptions about the distribution, assumptions about a linear or non-linear relationship, etc.

# (Non)parametric Correlation Measures


## Pearson's $\rho$

Linear relation between PCT\_\_WOMEN\_ON\_BOARD and EQY\_BETA?

```
cor(data$PCT_WOMEN_ON_BOARD, data$EQY_BETA, use="complete.obs")
```

```
## [1] -0.08745765
```

For confidence intervals and a significance test, use `cor.test()`.



# (Non)parametric Correlation Measures

## Spearman rank correlation

```
cor(data$PCT_WOMEN_ON_BOARD, data$EQY_BETA, use="complete.obs", method="spearman")
```

```
## [1] -0.08065364
```

## Kendall's $\tau$

```
cor(data$PCT_WOMEN_ON_BOARD, data$EQY_BETA, use="complete.obs", method="kendall")
```

```
## [1] -0.05475672
```

See [here](#) (Wikipedia) for comparisons between Spearman and Pearson correlations.

# Your turn

- Load the data in `sp500_data.csv` again
- Think about interesting questions and formulate hypotheses you could test with these data
- Test your hypothesis(-es) using appropriate tests
- Can you reject the  $H_0$ ?
- How do you interpret your results?
- Comment on upcoming issues regarding the statistical power, multiple hypotheses, normality assumptions, etc.